



Acceleration of three-dimensional Tokamak magnetohydrodynamical code with graphics processing unit and OpenACC heterogeneous parallel programming

H. W. Zhang, J. Zhu, Z. W. Ma, G. Y. Kan, X. Wang & W. Zhang

To cite this article: H. W. Zhang, J. Zhu, Z. W. Ma, G. Y. Kan, X. Wang & W. Zhang (2019) Acceleration of three-dimensional Tokamak magnetohydrodynamical code with graphics processing unit and OpenACC heterogeneous parallel programming, International Journal of Computational Fluid Dynamics, 33:10, 393-406, DOI: [10.1080/10618562.2019.1683167](https://doi.org/10.1080/10618562.2019.1683167)

To link to this article: <https://doi.org/10.1080/10618562.2019.1683167>



Published online: 30 Oct 2019.



Submit your article to this journal [↗](#)



Article views: 122



View related articles [↗](#)



View Crossmark data [↗](#)



Citing articles: 2 View citing articles [↗](#)



Acceleration of three-dimensional Tokamak magnetohydrodynamical code with graphics processing unit and OpenACC heterogeneous parallel programming

H. W. Zhang^a, J. Zhu^a, Z. W. Ma^a, G. Y. Kan^b, X. Wang^a and W. Zhang^a

^aInstitute for Fusion Theory and Simulation, Zhejiang University, Hangzhou, People's Republic of China; ^bState Key Laboratory of Simulation and Regulation of Water Cycle in River Basin, Research Center on Flood & Drought Disaster Reduction of the Ministry of Water Resources, China Institute of Water Resources and Hydropower Research, Beijing, People's Republic of China

ABSTRACT

In this paper, the OpenACC heterogeneous parallel programming model is successfully applied to modification and acceleration of the three-dimensional Tokamak magnetohydrodynamical code (CLT). Through combination of OpenACC and MPI technologies, CLT is further parallelised by using multiple-GPUs. Significant speedup ratios are achieved on NVIDIA TITAN Xp and TITAN V GPUs, respectively, with very few modifications of CLT. Furthermore, the validity of the double precision calculations on the above-mentioned two graphics cards has also been strictly verified with $m/n = 2/1$ resistive tearing mode instability in Tokamak.

ARTICLE HISTORY

Received 12 February 2019
Accepted 4 July 2019

KEYWORDS

OpenACC; MPI; GPU; magnetohydrodynamics (MHD)

1. Introduction

Magnetic confinement fusion is a method of using a magnetic field to confine a high-temperature fusion fuel, deuterium-tritium plasma, to generate thermonuclear fusion energy. There are different kinds of magnetically confined fusion devices in operation or under construction around the world, mainly Tokamaks, such as the DIII-D Tokamak (Evans et al. 2005) in the U.S., the Joint European Torus (JET) Tokamak (Liang et al. 2007) in Europe, the EAST Tokamak (Wan 2009) in China, the Wendelstein 7-X (W7-X) Stellarator (Renner et al. 2000) in Germany, the International Thermonuclear Experimental Reactor (ITER) under construction (Bécoulet et al. 2008), and the China Fusion Engineering Test Reactor (CFETR) under design (Song et al. 2014). Because phenomena observed in these devices are too complex to be studied analytically, computational simulation becomes a powerful tool to investigate their inside physical mechanisms.

Considering the restrict of computational capabilities of the single processor core, parallel programming plays a major role in the basic research and application fields of programme acceleration. Multiple parallel acceleration methods, such as MPI (Walker

and Dongarra 1996) and OpenMP (Dagum and Menon 1998), have been developed in past decades. In recent years, benefitting from the rapid performance improvement of the graphics processing unit (GPU), new parallel programming methods and tools, such as CUDA (Cook 2012), OpenCL (Kaeli et al. 2015), and OpenACC (Farber 2016), have been developed vigorously.

OpenACC as a user-driven directive-based performance-portable parallel programming model, is developed to simplify the parallel programming for scientists and engineers. Compared with the CUDA and OpenCL which require great efforts on code redevelopment, OpenACC has many advantages, such as satisfactory acceleration with very few modifications on an original source code and good compatibility with other devices, for example central processing unit (CPU). It has been successfully applied in some scientific and engineering codes, such as the flow code NeK5000 (Markidis et al. 2015), the computational electromagnetics code Nekton (Otten et al. 2016), the C++ flow solver ZFS (Kraus et al. 2014), the Rational Hybrid Monte Carlo (RHMC) QCD code (Gupta and Majumdar 2018), the Gyrokinetic Toroidal Code (GTC) (Wang et al. 2016), the space plasma

Particle-in-cell (PIC) code iPIC3D (Peng et al. 2015), the three dimensional pseudo-spectral compressible magnetohydrodynamic GPU code G-MHD3D (Mukherjee et al. 2018), the solar MHD code MAS (Caplan et al. 2018), and etc. The application perspective of the OpenACC technology in other scientific and engineering areas is very good and attractive.

CLT code (CLT is the abbreviation of Ci-Liu-Ti, which means magnetohydrodynamics in Chinese) is recently developed for the simulation study of magnetohydrodynamics (MHD) instabilities in toroidal devices, such as Tokamak (Wang and Ma 2015; Zhang, Ma, and Wang 2017). The outstanding issue is the poor parallel computational efficiency in accelerating the MHD code with MPI and OpenMP. With implementation of OpenACC, satisfactory acceleration of CLT with NVIDIA TITAN Xp and TITAN V GPUs is achieved. Compared with the speed of the MPI-parallelised CLT executed on 64-core CPUs (Intel® Xeon® Gold 6148F), we get about four times of acceleration with single TITAN V and double TITAN Xp GPUs. In addition, the simulation result on GPU is also verified in accuracy.

The paper is organised as follows: In Section 2, the MHD model and the main modules of CLT will be presented. In Section 3, the OpenACC implementation combined with MPI is given. And the acceleration performance of CLT with OpenACC is analysed in Section 4. The benchmark of simulation result from CLT with implementation of OpenACC is presented in Section 5. Finally, conclusions and discussion are given in Section 6.

2. Overview of CLT

CLT is an initial value full MHD code with toroidal geometries to study MHD instabilities in magnetically confined fusion devices, such as Tokamak. The code is written in FORTRAN 90 with double precision format and has grown to more than 20,000 code lines with hundreds of subroutines and functions. CLT has been successfully applied in studying the influence of toroidal rotation (Wang and Ma 2015), external driven current (Wang, Ma, and Zhang 2016; Zhang, Wang, and Ma 2017), and Hall effect (Zhang, Ma, and Wang 2017) on resistive tearing modes in Tokamaks. Its hybrid kinetic-magnetohydrodynamic version code, CLT-K, has also been developed to investigate dynamics of toroidal Alfvénic eigenmodes

in Tokamak (Zhu, Ma, and Wang 2016; Zhu et al. 2018).

2.1. Basic MHD equations

The MHD simulation is frequently used in magnetic confined fusion. In Tokamak, the time scale of an MHD instability is much longer than that of microscopic particle dynamics, and its spatial scale is much larger than the particle gyroradius. In this circumstance, the particle-in-cell method is failure due to its unbearable computation requirements, and the fluid approximation becomes quite effective and workable.

Apart from the continuity, momentum and energy equations widely used in computational fluid dynamics (CFD), MHD equations contain the Lorentz force in the momentum equation, Faraday's law, generalised Ohm's law, and Ampere's law neglecting displacement current. Because of the similarities between CFD and MHD, CFD solvers can be generalised to study MHD problems (Khodak 2015), and vice versa.

The full set of resistive MHD equations including dissipations can be written as follows:

$$\partial_t \rho = -\nabla \cdot (\rho \mathbf{v}) + \nabla \cdot [D \nabla (\rho - \rho_0)], \quad (1)$$

$$\partial_t p = -\mathbf{v} \cdot \nabla p - \Gamma p \nabla \cdot \mathbf{v} + \nabla \cdot [\kappa (p - p_0)], \quad (2)$$

$$\partial_t \mathbf{v} = -\mathbf{v} \cdot \nabla \mathbf{v} + (\mathbf{J} \times \mathbf{B} - \nabla p) / \rho + \nabla \cdot [\nu (\mathbf{v} - \mathbf{v}_0)], \quad (3)$$

$$\partial_t \mathbf{B} = -\nabla \times \mathbf{E}, \quad (4)$$

with

$$\mathbf{E} = -\mathbf{v} \times \mathbf{B} + \eta (\mathbf{J} - \mathbf{J}_0) + \frac{d_i}{\rho} (\mathbf{J} \times \mathbf{B} - \nabla p), \quad (5)$$

$$\mathbf{J} = \nabla \times \mathbf{B}, \quad (6)$$

where ρ , p , \mathbf{v} , \mathbf{B} , \mathbf{E} , and \mathbf{J} are the plasma density, thermal pressure, plasma velocity, magnetic field, electric field, and current density, respectively. $\Gamma (= 5/3)$ is the ratio of specific heat of plasma.

All variables are normalised as follows:

$$\begin{aligned} \mathbf{B}/B_m &\rightarrow \mathbf{B}, \mathbf{x}/a \rightarrow \mathbf{x}, \rho/\rho_m \rightarrow \rho, \mathbf{v}/v_A, t/\tau_a \\ &\rightarrow t, p/(B_m^2/\mu_0) \rightarrow p, \\ \mathbf{J}/(B_m/\mu_0 a) &\rightarrow \mathbf{J}, \mathbf{E}/(v_A B_m) \\ &\rightarrow \mathbf{E} \text{ and } \eta/(\mu_0 a^2/\tau_a) \rightarrow \eta, \end{aligned}$$

where $\tau_a = a/v_A$ is the Alfvénic time; $v_A = B_m/(\mu_0 \rho_m)^{1/2}$ is the Alfvénic speed; B_m and ρ_m are the

magnetic field and plasma density at the magnetic axis, respectively; a is the minor radius in the poloidal cross-section.

Constrained by the equilibrium conditions, the following equations should be satisfied:

$$\nabla \cdot (\rho_0 \mathbf{v}_0) = 0, \quad (7)$$

$$\mathbf{v}_0 \cdot \nabla p_0 + \Gamma p_0 \nabla \cdot \mathbf{v}_0 = 0, \quad (8)$$

$$\rho_0 \mathbf{v}_0 \cdot \nabla \mathbf{v}_0 = \mathbf{J}_0 \times \mathbf{B}_0 - \nabla p_0, \quad (9)$$

$$\nabla \times \mathbf{E}_0 = 0. \quad (10)$$

Substituting these equilibrium equations into Equations (1)–(4), the MHD equations can be rewritten as

$$\partial_t \rho = -\nabla \cdot (\rho \mathbf{v}_1 + \rho_1 \mathbf{v}_0) + \nabla \cdot [D \nabla (\rho - \rho_0)] \quad (11)$$

$$\begin{aligned} \partial_t p = & -\mathbf{v}_1 \cdot \nabla p - \mathbf{v}_0 \cdot \nabla p_1 \\ & - \Gamma (p \nabla \cdot \mathbf{v}_1 + p_1 \nabla \cdot \mathbf{v}_0) + \nabla \cdot [\kappa (p - p_0)], \end{aligned} \quad (12)$$

$$\begin{aligned} \partial_t \mathbf{v} = & -(\mathbf{v} \cdot \nabla \mathbf{v}_1 + \mathbf{v}_1 \cdot \nabla \mathbf{v}_0 + \rho_1 \mathbf{v}_0 \cdot \nabla \mathbf{v}_0 / \rho) \\ & + (\mathbf{J}_1 \times \mathbf{B} + \mathbf{J}_0 \times \mathbf{B}_1 - \nabla p_1) / \rho \\ & + \nabla \cdot [\nu (\mathbf{v} - \mathbf{v}_0)], \end{aligned} \quad (13)$$

$$\partial_t \mathbf{B} = -\nabla \times \mathbf{E}_1, \quad (14)$$

where the variables with subscript 0 represent equilibrium components and 1 for perturbation components, e.g. $\mathbf{v}_1 = \mathbf{v} - \mathbf{v}_0$. Thus, numerical errors from equilibrium can be minimised.

2.2. Coordinate systems and numerical schemes

In CLT, the cylindrical coordinate system (R, ϕ, Z) as shown in Figure 1 is adopted. The three-dimensional simulation boundary determined by the last closed magnetic surface and the $m/n = 3/1$ magnetic island for the EAST Tokamak from CLT are also plotted in Figure 1. In the magnetic confinement fusion device such as Tokamak, R , ϕ and Z indicate the major radius, toroidal, and up-down directions, respectively. One advantage of this coordinate system is that we can avoid the singularity near the $r = 0$ point that occurs in the toroidal coordinate (ψ, θ, ζ) . However, the boundary handling would be more difficult in the cylindrical coordinate. In the previous version of CLT, the fixed boundary condition at the last flux surface of plasma is assumed. Recently, the cut-cell method (Duan, Wang, and Zhong 2010) has been applied in




Figure 1. A three-dimensional perspective view of magnetic surface geometry for the EAST Tokamak. The three brown ribbons indicate $m/n = 3/1$ magnetic island from CLT. The pink surface represents the last closed magnetic surface while the green surface is the simulation boundary. The adopted cylindrical coordinate system (R, ϕ, Z) is also given.

CLT successfully, the details of the new boundary handling method will be introduced in another paper.

The uniform and rectangular grid is used in the $R - Z$ directions. For the numerical scheme, the 4th order finite difference method is employed in the R and Z directions, while in the ϕ direction, either 4th order finite difference or pseudo-spectrum method can be used. As for the time-advance, the 4th order Runge–Kutta scheme is applied.

For the parallelisation on CPU platform, the simulation domain can be divided into multiple blocks in each direction, and four-layer grids are used for message passing between every two neighbouring blocks according to the 4th order finite difference method in space. Thus, the increase of MPI CPU cores with a fixed mesh size will lead to rapid deterioration of the acceleration efficiency.

2.3. Main modules and basic execution flow chart of CLT

For each iterative step of CLT, the eight components of $\rho, p, \mathbf{v}, \mathbf{B}$ need to be updated. The main processes of CLT are shown in Figure 2. The modules with red blocks are executed on CPU, while the blue ones on GPU, and the green blocks mark out the main OpenACC directives in the whole procedure. The **INITIAL** module reads the equilibrium data that is calculated by the equilibrium solving code, such as NOVA (Cheng and Chance 1987) and EFIT (Lao et al.

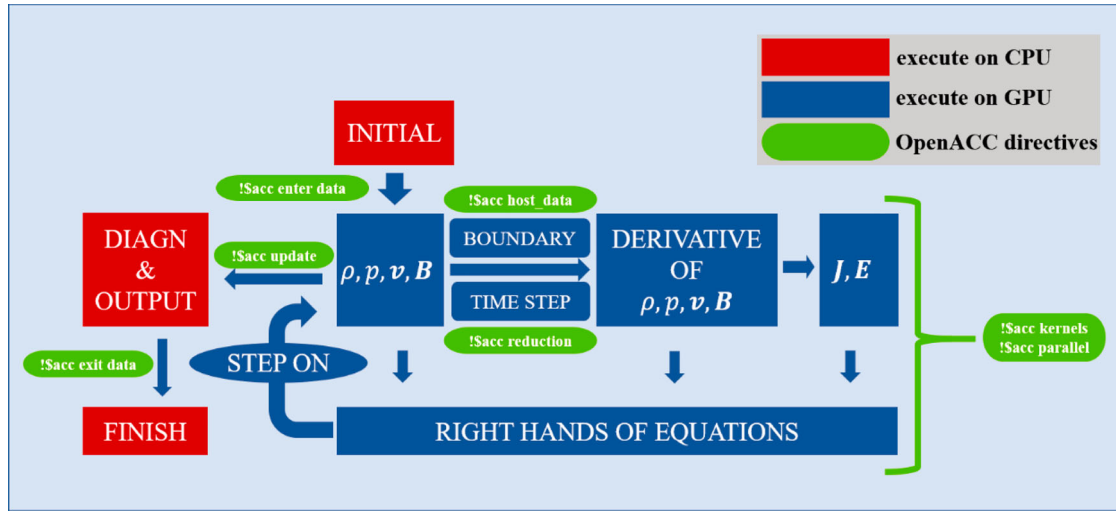


Figure 2. Main modules and OpenACC directives of CLT. The modules in the blue blocks are executed on GPU, and those in the red blocks are on CPU. The green blocks indicate the OpenACC directives applied in the corresponding positions.

1985). Then, the **DERIVATIVES OF ρ, p, v, B** module using the 4th order finite difference method is called to obtain the derivatives. With these derivatives, the values of the current **J** and electric field **E** are obtained. The derivatives, current **J** and electric field **E**, are substituted into the **RIGHT HANDS OF EQUATIONS**, and finally ρ, p, v, B are updated in **STEP ON** module with the 4th order Runge–Kutta scheme. Note that during the iteration the **BOUNDARY** module is applied in each step to solve the boundary conditions and deals with data exchange due to MPI parallelisation, while the **TIME STEP** module calculates the time increment based on the latest ρ, p, v, B to satisfy the Courant–Friedrichs–Lewy (CFL) condition in whole simulation domain, and some other modules such as **DIAGN** (recording data for time sequence analysis) and **OUTPUT** (recording data for global analysis) are called every few steps.

The execution time of CLT is mostly spent on the **RIGHT HANDS OF EQUATIONS**, **DERIVATIVES OF ρ, p, v, B** , and **J, E** modules. Therefore, a majority of OpenACC directives are added in these modules.

3. OpenACC and MPI implementations in CLT

3.1. OpenACC directives added to the serial version of CLT

The main idea of GPU-acceleration using OpenACC is to offload the calculation process from CPU to GPU. The data transfer speed between the memories of CPU and GPU is much slower than that of

GPU self-memory access. Thus, the first important operation is to copy all variables required in iterations from CPU memory to GPU memory. Therefore, after reading the equilibrium data in the **INITIAL** module, all necessary variables are copied into the GPU memory by using *\$sacc enter data copyin (variable name list)* directive and clause as shown in Figure 2. Accordingly, after finishing all iterations, the *\$sacc exit data delete (variable name list)* directive and clause are also used before the programme finishes to free the GPU memory. There are also other methods playing the same role, such as combination of *!\$sacc declare create (variable name list)* and *!\$sacc update device (variable name list)*. The main advantage of *enter data* and *declare create* construct is that the data lifetime in the GPU memory can be extended among different functions and subroutines. With this feature, the calculation on the GPU can be continuously proceeded without stopping for data transfer between CPU and GPU.

A demonstrative example for the *enter data* constructs in CLT is given in Figure 3. For simplicity, only three representative variables (*xx, yy, zz*) are listed in the OpenACC *copyin* clause. During the main iteration, CLT calls the **DIAGN** and **OUTPUT** modules every few steps to record data into hard disk by CPU. Because the **DIAGN** module must use some external math libraries, such as MKL (Wang et al. 2014), this module is still executed on CPU due to its difficulty for parallelisation with OpenACC. Therefore, just before calling the **DIAGN** and **OUTPUT** modules, the directive *!\$sacc update host (variable name list)* is used to


```

PROGRAM CLT
! INITIALIZE
CALL INITIA
!$acc enter data copyin(xx, yy, zz)
! MAIN ITERATION PART
DO NSTEP = 1, NSTEP
    CALL STEPON
    IF(MOD(NSTEP,NDIAGN) == 0) THEN
        !$acc update host(xx, yy, zz)
        CALL DIAGN
    ENDIF
    IF(MOD(NSTEP,NOUTPUT) == 0) THEN
        !$acc update host(xx, yy, zz)
        CALL OUTPUT
    ENDIF
ENDDO
!$acc exit data delete(xx, yy, zz)
END PROGRAM CLT

```

Figure 3. An example for 'enter data' and 'update' constructs in CLT.

copy the latest data from GPU back to CPU. The time wasted on updating operation is negligible because the number of calls of these two modules is rather limited.

Since CLT is a three-dimensional MHD code, the majority of loops contains three levels. To parallelise the loops, two constructs of *kernels* and *parallel* are frequently used. A representative example for such parallelisation in the **TIME STEP** module is given in Figure 4. The *present* clause following *parallel* construct directive indicates that the variables inside the parentheses already exist on GPU, which can avoid the needless implicit variable copy or create operations on GPU. The *loop* clause after *parallel* directive should be placed just before the loop body, which instructs the compiler to perform parallelisation for the closely followed loop. As for the **TIME STEP** module, the reduction for the minimum time step is used by adding *reduction(min:dt1)*, because the time step used in time advance must satisfy the CFL condition in each grid. Other reduction operations are also supported in OpenACC with similar syntaxes (Farber 2016). The *independent* clause instructs the compiler that the calculation on each grid has no dependency and therefore can be parallelised directly. The final clause in Figure 4 is *collapse(3)*. Collapsing loops means that, for example, three loops of trip counts NX, NZ, and NY will automatically combine into a single loop with a trip total count of NX*NZ*NY, which greatly increases the parallelisation efficiency. Note that the *parallel* constructs

```

SUBROUTINE SETDT
USE DECLARE
INCLUDE 'MPIF.H'
DT1=100.d0
!$acc parallel present(x, xx, yy, zz, gdt_p_ep)
!$acc loop reduction(min:dt1) independent collapse(3)
DO JY = IY_FIRST + 2, IY_LAST - 2
DO JZ = IZ_FIRST + 2, IZ_LAST - 2
DO JX = IX_FIRST + 2, IX_LAST - 2
    IF(GDTP_EP(JX,JZ,1).NE.4) THEN
        VX=X(JX,JZ,JY,3)
        VY=X(JX,JZ,JY,4)
        VZ=X(JX,JZ,JY,5)
        VA2=(X(JX,JZ,JY,6)**2+X(JX,JZ,JY,7)**2 &
            +X(JX,JZ,JY,8)**2)/X(JX,JZ,JY,1)
        CS2=GAMMA*X(JX,JZ,JY,2)/X(JX,JZ,JY,1)

        VPX=DABS(VX)+SQRT(DABS(CS2+VA2))
        VPY=DABS(VY)+SQRT(DABS(CS2+VA2))
        VPZ=DABS(VZ)+SQRT(DABS(CS2+VA2))

        DTX=DABS(XX(JX)-XX(JX-1))/(VPX/CFL)
        DTZ=DABS(ZZ(JZ)-ZZ(JZ-1))/(VPZ/CFL)
        DTY=DABS(XX(JX)*(YY(JY)-YY(JY-1)))/(VPY/CFL)

        DT2=DMIN1(DTX,DTZ)
        DT3=DMIN1(DTY,DT2)
        DT1=DMIN1(DT1,DT3)
    ENDIF
ENDDO
ENDDO
ENDDO
!$acc end parallel
RETURN
END SUBROUTINE SETDT

```

Figure 4. An example for the parallelisation in the **TIME STEP** module using 'parallel' constructs.

used here can also be replaced by *kernels* construct, the major difference of these two constructs is that *kernels* gives the compiler more freedom to find and map parallelism according to the requirements of the target accelerator, while the *parallel* construct is more explicit, and requires more analysis by the programmer.

Another important OpenACC directive used in CLT is the *!\$acc routine* directive for the procedure call. Figure 5 gives the solution when a subroutine or function is called inside an accelerated region. The subroutine *interp1d2l* interpolates data between different mesh grids, and is commonly used in the accelerated region of **BOUNDARY** module. The *!\$acc routine seq* directive instructs the compiler to generate a device version for this subroutine on GPU so that it can be called in the accelerated region. An interface for this child subroutine is required inside the parent as given in Figure 5. Then the subroutine can be called inside the OpenACC accelerated region directly.

```

SUBROUTINE INTERP1D2L(X1, X2, X3, Y1, Y2, Y3, Y, ANS)
!$acc routine seq
REAL*8 X1,X2,X3,Y1,Y2,Y3,Y,ANS
REAL*8 D1,D2,D3
D1 = (Y1-Y2)*(Y1-Y3)
D2 = (Y2-Y3)*(Y2-Y1)
D3 = (Y3-Y1)*(Y3-Y2)
ANS = X1*(Y-Y2)*(Y-Y3)/D1 &
      + X2*(Y-Y3)*(Y-Y1)/D2 &
      + X3*(Y-Y1)*(Y-Y2)/D3

RETURN
END SUBROUTINE INTERP1D2L
!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
SUBROUTINE BOUNDARY
USE DECLARE
IMPLICIT NONE
INCLUDE 'MPIF.H'

! INTERFACE FOR ACCELERATED SUBROUTINE
INTERFACE
  SUBROUTINE INTERP1D2L(X1, X2, X3, Y1, Y2, Y3, Y, ANS)
    !$acc routine seq
    REAL*8 X1,X2,X3,Y1,Y2,Y3,Y,ANS
    REAL*8 D1,D2,D3,TMP_ADD
  END SUBROUTINE
END INTERFACE

! OpenACC ACCELERATED REGION
.....
CALL INTERP1D2L(X1, X2, X3, Y1, Y2, Y3, Y, ANS)
.....
! END OF OpenACC ACCELERATED REGION

END SUBROUTINE BOUNDARY

```

Figure 5. An example for the procedure call (subroutine `interp1d2l`) inside OpenACC accelerated region.

3.2. Combination of OpenACC and MPI technologies for parallelisation on multiple GPUs

CLT was developed with MPI at first for parallelisation in clusters. For a case with (256, 256, 32) grids in the (R, Z, ϕ) directions, the GPU memory usage in a graphics card is about 2.25GB. Thus, the memory in a single graphics card such as TITAN Xp or TITAN V (with 12GB of GPU memory) is more than enough. However, for a high-resolution simulation, a size of (1024, 1024, 128) grids or even larger should be chosen. In such scenarios, more than 64 times memory (over 144GB) will be required on the GPU. Naturally, the multiple GPUs parallelisation becomes indispensable. Based on the original MPI in CLT, we also added the directives for directly communication between GPUs without transferring data back to CPU.

As shown in Figure 6, firstly, we add the `!$acc set device_num` (*rank of GPU*) directive just after the `mpi_comm_size` and `mpi_comm_rank` calls. The value of *rank of GPU* can be defined as functions of the

```

PROGRAM CLT
.....
! INITIATE MPI
CALL MPI_INIT(IERROR)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, NSIZE, IERROR)
CALL MPI_COMM_RANK(MPI_COMM_WORLD, NRANK, IERROR)
!$acc set device_num(0)
.....
!MPI SEND-----
!$acc host_data use_device(wfx1) if_present
CALL MPI_SEND(WFX1, MYZ8, MPI_DOUBLE_PRECISION, NRANK+1, 0, &
  MPI_COMM_WORLD, IERROR)
!$acc end host_data
.....
!MPI RECEIVE-----
!$acc host_data use_device(wfx1) if_present
CALL MPI_RECV(WFX1, MYZ8, MPI_DOUBLE_PRECISION, NRANK-1, 0, &
  MPI_COMM_WORLD, IERROR)
!$acc end host_data
.....
END PROGRAM CLT

```

Figure 6. An example for MPI called on GPU by using '`host_data use_device`' construct.

process rank in MPI parallelisation so that each MPI rank can be assigned to a specified GPU for the load balance. With this *rank of GPU*, the communication between different GPUs becomes feasible. Then, a specific example with OpenACC directives in `mpi_send` and `mpi_recv` API is given. The `!$acc host_data use_device (variable name list) if_present` instructs the compiler to use the address of the specified variables on GPU if the variables are present, then the `mpi_send` or `mpi_recv` API is executed by accessing the data from the GPU address of variable rather than the CPU. The adoption of the `host_data use_device` construct can avoid the unnecessary data transfer between CPU and GPU before MPI calls. For the purpose of comparison, an equivalent example is given in Figure 7, where the MPI still accesses the data from the CPU address for specified variables. However, in order to update the data inside the CPU and GPU, another pair of `!$acc update host (variable name list)` and `!$acc update device (variable name list)` are added before the `mpi_send` and after the `mpi_recv`, respectively. The functions of the methods used in Figures 6 and 7 are the same, but the later one will spend additional time on data updating between CPU and GPU.

4. Acceleration performance analysis

The performance tests are carried out on two platforms for CPU and GPU version of CLT, respectively.

```

PROGRAM CLT
.....
! INITIATE MPI
CALL MPI_INIT(IERROR)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD,NSIZE,IERROR)
CALL MPI_COMM_RANK(MPI_COMM_WORLD,NRANK,IERROR)
!$acc set device_num(0)
.....
!MPI SEND-----
!$acc update host(wfx1)
CALL MPI_SEND(WFX1,MYZ8,MPI_DOUBLE_PRECISION, NRANK+1, 0, &
MPI_COMM_WORLD, IERROR)
.....
!MPI RECEIVE-----
CALL MPI_RECV(WFX1,MYZ8,MPI_DOUBLE_PRECISION, NRANK+1, 0, &
MPI_COMM_WORLD, IERROR)
!$acc update device(wfx1)
.....
END PROGRAM CLT

```

Figure 7. An example for MPI called on CPU by using ‘update’ construct.

The detailed information for each platform is given in Table 1, including the hardware configurations, compiler information, compiler options, and etc. The CPU version of CLT is executed on KYLIN-2 cluster with Intel® Fortran compiler included in the Intel® Parallel Studio XE, while the GPU version is executed on x3860-x6 workstation with 2 TITAN Xps and 2 TITAN Vs installed, the compiler is pgfortran 18.10 included in the PGI Compilers and Tools. Note that the compiler options are different for TITAN Xp and TITAN V due to their different computational capabilities. And the OpenMPI on GPU platforms is supported by the PGI Compilers and Tools, which is used in the combination of OpenACC and MPI technologies for parallelisation on multiple GPUs as discussed above.

4.1. Performance on single GPU card

For the case with grids as (256, 256, 16) in R , Z , and ϕ , the comparison for execution time (with 20,000 steps) is carried out on different platforms. Some simulation results are presented in Figure 8. The MPI block division for the CPU version of CLT is carried out uniformly in the R and Z directions, while no block division is applied in the GPU version because we only used one card for acceleration. Note that for the case executing on the CPU, the ‘O3’ compiler option is turned on so as to obtain the best performance optimisation.

As shown in Figure 8, the performance of the code based on MPI executed on CPU is still acceptable when total 128 cores (4 nodes) are used, with the parallel

Table 1. Detailed information of two platforms used in the acceleration performance tests.

Platform	CPU configuration	Memory configuration	GPU configuration	Fortran compiler version	MPI version	Compiler options
KYLIN-2 cluster CentOS Linux release 7.4.1708	138 nodes, with 2 Intel® Xeon® Gold 6148F CPUs (i.e. 40 cores) in each node	92GB in each node	N/A	Intel(R) Fortran Intel(R) 64 Compiler XE for applications running on Intel(R) 64, Version 15.0.3.187 Build 20150407	mpiifort for the Intel(R) MPI Library 5.0 Update 3 for Linux*	-O3, -mcmmodel = large
x3860-x6 workstation Red Hat Enterprise Linux Server release 6.8	1 node, with 16 Intel Xeon E7-8867 v4 CPUs (i.e. 144 cores in total)	1024GB in total	2 TITAN Xps & 2 TITAN Vs	pgfortran 18.10-0 64-bit target on x86-64 Linux -tp haswellPGI Compilers and Tools	mpif90, OpenMPI 2.1.2	-acc, -mcmmodel = medium, -Mlarge_arrays, -ta = tesla:cc70 (for TITAN V) & -ta = tesla:cc60 (for TITAN Xp)

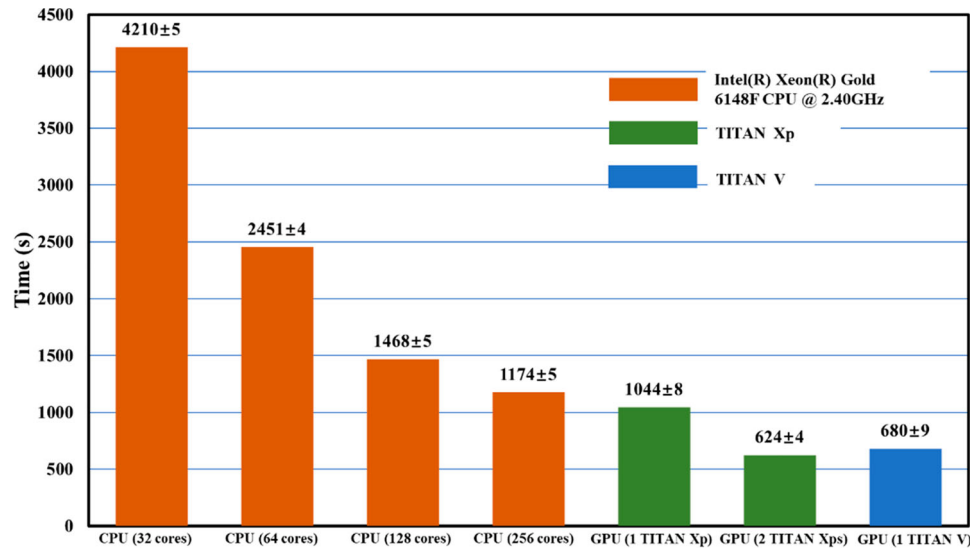


Figure 8. Comparison for execution time on different platforms (CPU coloured by orange, TITAN Xp by green, and TITAN V by blue).

acceleration efficiency at 72%. However, the efficiency drops quickly to 45% when the number of CPU cores is increased to 256 (7 nodes). The reason for the decline of the parallel acceleration efficiency with increase of CPU cores is mainly due to the growth of the overlap area between different MPI blocks with 4th order finite difference scheme in space. Therefore, the speedup for an MHD code with a fixed mesh is quite difficult with traditional MPI. Fortunately, the application of OpenACC based on GPU on our CLT can achieve a quite satisfactory acceleration efficiency. As the results shown in Figure 8 with the green and blue bars for TITAN Xp and TITAN V, respectively, the elapsed time on single TITAN Xp almost equals to that of 256 Intel® Xeon® Gold 6148F CPU cores. What's more, the speed of TITAN V is even 54% faster than that of TITAN Xp. Compared with our frequently-used MPI set of 64 cores, the speedup of TITAN V is about 3.6 times. In addition to the speed performance, another advantage for the GPU acceleration is the much lower price for constructing GPU workstations or servers than building a cluster with CPU nodes. The cost of a single TITAN V is comparable with one Intel® Xeon® Gold 6148F CPU (20 cores in each). Thus, the cost performance of the former is about ten times better than the latter for our CLT.

Another important feature used in OpenACC is the combination of three levels parallelisation as *gang*, *work*, and *vector*. For the case shown in Figure 8, we found that the speed of CLT on the specified mesh is most sensitive to the number of *gang*: enough *gang*

parallelisation (more than 512) leads to the best performance on GPU, while too few gangs (less than 100) results in slowing down of the code compared with the MPI case with 32 cores. The default set of the compiler without adding any artificially specified configurations also leads to the best performance of the acceleration as shown in Figure 8. The numbers of *gang*, *worker*, and *vector* will affect the usage of GPU memory cache and can influence the speed greatly. For an MHD code accelerated with OpenACC, adjusting these configurations is necessary to obtain the best acceleration performance, and the best set is also quite different for different codes, devices, and problem sizes.

4.2. Performance on double GPU cards combined with MPI

For the lack of cluster with multiple GPU nodes, the speed test for the OpenACC combined with MPI has only been performed on up to two separate GPU cards inside a workstation. Note that the computation and communication in CLT is separated, that is, the synchronous communication method is adopted.

Firstly, a simple test for MPI-GPU combined acceleration is carried out with 4 MPI ranks on single TITAN Xp with the two implementation methods mentioned in Figure 6 and 7 (by setting 4 MPI threads in a same GPU rank). The first method (as shown in Figure 6, with direct communication between GPUs) leads to only about 15% slowing down compared with time showed by the green bar in Figure 8,

while the other method (as shown in Figure 7, using *update* clause and communication between CPU MPI threads) leads to about 105% slowing down. Therefore, the support of direct MPI communication on GPU makes the multiple GPU acceleration for large size problems feasible, while the data exchange between GPU and CPU memories will intolerably slow down the execution speed.

Secondly, a corresponding test for MPI-GPU combined acceleration is performed on two TITAN Xp GPUs. The method of direct communication between GPUs as shown in Figure 6 is chosen. Two MPI processes in total are carried out with double TITAN Xp GPUs, that is, each GPU executes one MPI process, and the computational domain is equally divided into two blocks in the R direction. As the result shown in Figure 8, the final execution time is about 624 s, which is even faster than that on single TITAN V. Compared with the performance of single TITAN Xp, the parallel efficiency is about 84%. The communication time is relatively small compared with the computation time, though the communication in CLT is synchronous. Therefore, the time cost on data interaction through Peripheral Component Interface Express (PCIe) is negligible and the parallelisation with multiple GPUs is workable, which is necessary and indispensable for simulating larger size problem in future research.




Figure 9. Execution times on single or double TITAN Xp GPUs for different mesh size in R direction (N_R equals to 32, 64, 128, 256, and 512, respectively, while N_Z and N_ϕ are fixed to be 256 and 16). Red circles indicate the execution times on single TITAN Xp, and blue squares indicate the times on double TITAN Xps. The communication times are plotted with black stars. Linear fitting lines for each case are given. The parallel efficiencies are also plotted with green triangles.

In order to further examine the parallel scalability of CLT, the execution times for different mesh size N_R in the R direction in single and double TITAN Xps are studied and given in Figure 9 (N_R equals to 32, 64, 128, 256 and 512, respectively, while N_Z and N_ϕ are fixed to be 256 and 16). The computational domain divisions for all cases with double TITAN XPs are the same as above. Scaling laws with linear fitting lines are plotted in Figure 9. The additional communication time resulted from MPI parallelisation is estimated by the difference of the execution time on double GPUs and half of that on single GPU. It is found that the communication time are weakly dependent on the value of N_R though the sizes of communication area are controlled to be identical. With these scaling laws, the parallel efficiency P for a specified MPI domain decomposition with mesh size of (N_R , 256, 16) in R , Z , and ϕ can be estimated with

$$P = \frac{3.56N_R + 100.00}{\frac{0.13N_R + 63.92}{256} \times L_{Communication}^{Maximum} + (3.56N_R + 100.00)/N_{GPUs}} \times \frac{1}{N_{GPUs}} \quad (15)$$

where the numerator estimates the execution time on single TITAN Xp, while the denominator estimates that on N_{GPUs} TITAN Xps. The first term in denominator estimates the communication time (the communication time is assumed to be proportional to the maximum length of grids $L_{Communication}^{Maximum}$ in all processors used for communication, for example, $L_{Communication}^{Maximum}$ equals to 256 in all cases of Figure 9). And the second term in denominator estimates the computation time, which is assumed to be inversely proportional to the number of GPUs. With Equation (15), the estimated parallel efficiencies for four cases in Figure 9 are 61%, 69%, 78%, 84% and 88% respectively for N_R equals to 32, 64, 128, 256 and 512, while the actual efficiency values are 55%, 68%, 81%, 84% and 88%. Consequently, the estimation formula is relatively accurate though there is still some error. As for the mesh size of (256, 256, 16) in R , Z , and ϕ used in simulation, the parallel efficiency for 4 GPUs parallelisation is estimated to be 72%, and that value drops to a low value of 51% for 8 GPUs. Therefore, the optimal number of GPUs for mesh size of (256, 256, 16) is expected to be 4 on which hopefully about threefold acceleration could be achieved compared with single GPU.

4.3. Performance study for the double precision computational capability

Another issue of the acceleration performance of the GPU-OpenACC must be discussed here is the computational capability of the double precision floating points. The theoretical double precision capability of TITAN Xp is up to 0.38 TFLOPS, which is only comparable with one Intel[®] Xeon[®] Gold 6148F CPU with 20 cores. However, the execution speed of CLT on TITAN Xp is about 4 times faster than that on the Gold 6148F CPU with 32 cores parallelisation. Although the double precision computational capability of TITAN V is about 6.90 TFLOPS that is more than 10 times of TITAN Xp, the real execution speed on TITAN V only achieves less than two times on TITAN Xp. Therefore, the parallelisation performance of the code on different platforms does not entirely depend on the theoretical computational capability. Other factors such as the memory accessing time, partition, bandwidth, I/O situation, compiler optimisation options, and the network conditions will also affect the execution speed of the code. To study this unexpected result, a simple demo in double precision as shown in Figure 10 is written to investigate the influence of the computational intensity on acceleration efficiency. One parameter `N_CYCLE` in this demo is used to control the computational intensity. When `N_CYCLE` is set to a small value, for example, 1 (which is similar to the situation of CLT), the demo code is a memory intensive task, which means memory access for data read & write is required every few steps. However, if the `N_CYCLE` value becomes large, such as 100, the demo code becomes a computationally intensive task. Because after each memory access, the data can be locally used and updated for `N_CYCLE` times without pausing. By scanning the value of `N_CYCLE`, the time costs for the iterations and the bandwidth utilisation of the device memory (only for GPU cases) are compared for Gold 6148F CPU, TITAN Xp, and TITAN V, respectively. The NVIDIA (R) CUDA command line profiler `nvprof` is used to obtain the execution time and bandwidth utilisation of the kernel in Figure 10. The ‘`dram_read_throughput`’, ‘`dram_write_throughput`’, and ‘`dram_utilization`’ options in metrics mode of `nvprof` provide the device memory read & write throughput and the utilisation level of the device memory relative to the peak utilisation on a scale of low (0) to max (10). The detailed results are presented in Table 2.

```

PROGRAM DEMO
IMPLICIT NONE
REAL*8, DIMENSION(256,256,16,8) :: A, B, C, D
REAL*8 :: TIME_START, TIME_END, TIME_COST
INTEGER :: JX, JZ, JY, M, I_CYCLE, N_CYCLE
!$acc set device_num(0)

N_CYCLE = 1 ! N_CYCLE CONTROLS THE NUMBER OF CIRCULATIONS
A(:, :, :, :) = 1.D-6
B(:, :, :, :) = 1.D-6
C(:, :, :, :) = 1.D-6
D(:, :, :, :) = 1.D-6
!$acc enter data copyin(A,B,C,D)

CALL CPU_TIME(TIME_START)
!$acc kernels default(present)
!$acc loop independent collapse(4)
DO M=1,8
DO JY=2,15
DO JZ=2,255
DO JX=2,255
DO I_CYCLE=1,N_CYCLE
A(JX,JZ,JY,M) = A(JX,JZ,JY,M) + A(JX+1,JZ,JY,M) - A(JX-1,JZ,JY,M) + &
B(JX,JZ+1,JY,M) - B(JX,JZ-1,JY,M) + &
C(JX,JZ,JY+1,M) - C(JX,JZ,JY-1,M) + &
DSIN(D(JX,JZ,JY,M)) + DCOS(D(JX,JZ,JY,M))

ENDDO
ENDDO
ENDDO
ENDDO
ENDDO
!$acc end kernels
CALL CPU_TIME(TIME_END)
TIME_COST = TIME_END - TIME_START
PRINT*, 'N_CYCLE = ', N_CYCLE, 'TIME COST = ', TIME_COST

!$acc exit data delete(A,B,C,D)
END PROGRAM DEMO

```

Figure 10. A demo code to investigate the influence of computing intensity on acceleration efficiency. The demo code is memory intensive for small `N_CYCLE` value while the demo code tends to become compute intensive for large `N_CYCLE` value.

As shown in Table 2, when the demo code is a memory intensive task, i.e. `N_CYCLE` = 1, the speed up ratio of TITAN V and TITAN Xp is close to but slightly greater than that of CLT, with a value of about 2.103. In the meanwhile, the ratio of sustained bandwidth to theoretical peak bandwidth of TITAN V is about 0.853, and the maximal utilisation level of device memory has been reached according to the information provided by `nvprof`. In this case, compared with TITAN Xp, TITAN V does not show much advantage in spite of its powerful computational capability of double precision due to the bandwidth bound. However, when `N_CYCLE` is increased to a larger value, for example, `N_CYCLE` = 100, the code becomes computationally intensive. Correspondingly, the ratio of sustained bandwidth to theoretical peak bandwidth of TITAN V is only about 0.056, and the device memory utilisation is in the low level. As a result of the low

Table 2. Time costs comparisons for the demo code in Figure 10 executed at different platforms by scanning the value of N_CYCLE from 1 to 100. Note that the execution time for the kernel in the GPU case is obtained with nvprof.

TITAN Xp (Theoretical peak bandwidth: 547.680GB/s)					
N_CYCLE	Time (sec)	Sustained bandwidth (GB/s)	Bandwidth ratio (sustained/peak)	DRAM utilisation (scale from 0 to 10)	CPU Time (sec)
1	$(1.251 \pm 0.011)E-03$	264.549 ± 2.378	0.483 ± 0.004	Mid (6)	0.107 ± 0.004
2	$(2.466 \pm 0.022)E-03$	133.297 ± 1.200	0.243 ± 0.002	Low (3)	0.135 ± 0.005
5	$(6.098 \pm 0.062)E-03$	53.206 ± 0.008	$(9.715 \pm 0.002)E-02$	Low (2)	0.301 ± 0.005
10	$(1.225 \pm 0.001)E-02$	26.655 ± 0.005	$(4.867 \pm 0.001)E-02$	Low (1)	0.595 ± 0.010
100	$(1.196 \pm 0.051)E-01$	3.130 ± 0.282	$(5.716 \pm 0.515)E-03$	Low (1)	6.956 ± 0.047
TITAN V (Theoretical peak bandwidth: 652.800GB/s)					
N_CYCLE	Time (sec)	Sustained bandwidth (GB/s)	Bandwidth ratio (sustained/peak)	DRAM utilisation (scale from 0 to 10)	Speedup ratio (TITAN V/Xp)
1	$(5.950 \pm 0.009)E-04$	556.768 ± 0.504	0.853 ± 0.001	Max (10)	2.103 ± 0.016
2	$(5.977 \pm 0.012)E-04$	545.165 ± 2.382	0.835 ± 0.004	High (9)	4.125 ± 0.044
5	$(6.420 \pm 0.018)E-04$	506.097 ± 0.599	0.775 ± 0.001	High (9)	9.500 ± 0.104
10	$(1.025 \pm 0.001)E-03$	317.876 ± 0.481	0.487 ± 0.001	Mid (6)	11.952 ± 0.014
100	$(9.394 \pm 0.007)E-03$	36.665 ± 0.053	$(5.617 \pm 0.008)E-02$	Low (1)	12.734 ± 0.547

bandwidth utilisation, the speed of TITAN V is more than ten times faster than that of TITAN Xp, which is consistent with the difference of double precision computational capabilities between two GPU cards. Similar analysis for CLT has been done with nvprof, and the result demonstrates that CLT is bandwidth limited on TITAN V. Therefore, the bandwidth analysis of Table 2 indicates that TITAN V has a great acceleration potential for CLT by focusing on the data locality and memory access, which requires more efforts on the deep optimisation of the code.

5. Results of benchmarking

To verify of the results calculated by the OpenACC version CLT on the GPU, the simulations of $m/n = 2/1$ resistive tearing mode in Tokamak are carried out both on the CPU-MPI platform (Intel® Xeon® Gold 6148F CPU @ 2.40 GHz, 32 cores, ECC memory support) and the GPU-OpenACC platform (TITAN V or TITAN Xp, non-ECC memory support, using one GPU for each test).

The initial equilibrium is calculated by the NOVA code (Cheng and Chance 1987) with the initial safety factor q and pressure p profile as shown in Figure 11. The $m/n = 2/1$ tearing mode is most unstable for this equilibrium. The grid set is (256, 256, 16) in R , Z , and ϕ . The resistivity η is chosen to be 1×10^{-5} .

Temporal evolutions for the kinetic energy of the system given in Figure 12 are obtained from the CPU-MPI and GPU-OpenACC versions of CLT. The kinetic energies of the two cases are exactly the same during the whole 400,000 simulation steps, except for some




Figure 11. Initial equilibrium safety factor q profile and pressure p profile for simulation of $m/n = 2/1$ resistive tearing mode.

negligible differences on the 14th decimal place of the number, which is almost the machine error limitation of the double precision.

The mode structures from these two calculations are also identical, as the result of the CPU-MPI case demonstrated in Figure 13(a), and the contour plot for the difference of the mode structures by CPU and GPU is also given in Figure 13(b). The mode structure of the toroidal electric field E_ϕ in Figure 13(a) is clearly the $m/n = 2/1$ tearing mode, the result calculated by GPU-OpenACC is omitted due to the indistinguishability. The difference of the results between CPU-MPI and GPU-OpenACC as shown in Figure 13(b) is less than 10^{-14} , which is consistent with the situation of kinetic energy. As for the result on TITAN Xp, it is




Figure 12. Temporal evolutions for the kinetic energy of the system for two cases executed on CPU-MPI platform (blue line) and GPU-OpenACC platform (red circles), respectively.

just the same as that on TITAN V, thus it will not be repeated here.

Therefore, the consistency for the $m/n = 2/1$ tearing mode on both platforms confirms the reliability of GPU-OpenACC for the double precision simulation of CLT. The stability is also verified with dozens of simulations with each over 400,000 steps. The lack of ECC memory on TITAN Xp and TITAN V does not influence the correctness of the CLT results so far.

6. Discussion and conclusion

The successful application of OpenACC (Kan, Zhang, et al. 2018; Kan, He, Li, et al. 2019; Kan, Lei, et al. 2017; Kan, Liang, et al. 2016; Kan et al. 2018; Kan, He, Ding, et al. 2017) in CLT leads to great improvement in computational efficiency for studying the MHD instability

in Tokamak. The migration of CLT from the CPU-MPI platform to GPU-OpenACC platform is relatively easy compared with completely rewriting the code into CUDA or OpenCL. Only about 500 lines OpenACC directives have been added into CLT, which is quite few of lines compared with the total of more than 20,000 lines. Some minor changes have been done on the code, such as merging some small subroutines and adjusting the order of cyclic indices to obtain the best acceleration performance.

Compared with the speed of the same CLT case executed on CPU with 64 cores MPI-parallelisation, about four times of acceleration has been achieved on a single TITAN V and double TITAN Xp GPUs. The combination of MPI and OpenACC makes the multiple GPU acceleration to become feasible by simply using MPI with the data addressed on GPU memory.

It is found that the speedup does not just depend on the theoretical computational capabilities of the CPU and GPU. Only less than double speed of TITAN Xp is achieved on TITAN V for CLT, while the theoretical computational capability of the latter in double precision is ten times more than that of the former. By increasing the compute intensive of the demo code, the speed of TITAN V is more than ten times faster than that of TITAN Xp. Therefore, greater potential on TITAN V is observed if more optimizations on CLT, such as increase of data locality, are carried on. However, the speed of the code and the time spent on code modification and optimisation is a trade-off. Because the speedup performance on TITAN V is already satisfactory considering its low price, and to keep the integrity and readability of code structure, acceleration work on CLT is paused at this stage.




Figure 13. (a) The $m/n = 2/1$ tearing mode structure of the toroidal electric field E_ϕ and (b) the difference of E_ϕ from CPU-MPI platform and GPU-OpenACC platform.

Most important of all, the benchmarking for the results calculated by GPU are done by comparing that of traditional CPU platforms. The results from both platforms show exactly identical after 400,000 steps for dozens of runs. The double-precision operations on TITAN Xp and TITAN V can be fully trustable for CLT.

In addition, the migration to OpenACC for kinetic part of the hybrid kinetic-magnetohydrodynamic version, CLT-K, is still under way, which requires more modifications on the code structure and adjustment on the combination of different OpenACC directives. The experience for this work on CLT-K will be introduced in a future paper if the acceleration performance is noteworthy.

As magnetic confinement fusion is approaching the era of burning plasma, for instance, the under-construction project ‘International Thermonuclear Experimental Reactor (ITER)’ and the under-design project ‘China Fusion Engineering Test Reactor (CFETR)’, simulation studies become much more complicate and large scale which requires that a simulation code has to be efficient. With the GPU acceleration, the simulation period for a typical MHD case reduces from days to hours, which helps us to understand the physics in Tokamak greatly.

Migration CLT from a CPU machine to GPU with OpenACC can be taken as a reference for both CFD and MHD codes. The GPU-acceleration can be applied in more fields other than fusion research, such as the geo/astrophysics, bioscience, environics.

Acknowledgments

The authors would like to acknowledge helpful suggestions given by professor C. Yang, and the Sunway TaihuLight Supercomputer Team at National Supercomputing Center, Wuxi, China.

Disclosure statement

No potential conflict of interest was reported by the authors.

Funding

This work is supported by the Special Project on High-performance Computing under the National Key R&D Program of China [grant number 2016YFB0200603], the National Natural Science Foundation of China under [grant numbers 11775188 and 11835010], Fundamental Research Fund for Chinese Central Universities, Beijing Natural Science Foundation [grant number 8184094], IWHR Research & Development Support Program [grant number JZ0145B022018].

ORCID

H. W. Zhang  <http://orcid.org/0000-0003-1342-5756>

References

- Bécoulet, M., Eric Nardon, G. Huysmans, W. Zwingmann, P. Thomas, M. Lipa, R. Moyer, T. Evans, V. Chuyanov, and Y. Gribov. 2008. “Numerical Study of the Resonant Magnetic Perturbations for Type I Edge Localized Modes Control in ITER.” *Nuclear Fusion* 48 (2): 024003.
- Caplan, Ronald M, Jon A Linker, Zoran Mikić, Cooper Downs, Tibor Török, and V. S. Titov. 2018. “GPU Acceleration of an Established Solar MHD Code Using OpenACC.” arXiv preprint arXiv:1811.02605.
- Cheng, C. Z., and M. S. Chance. 1987. “NOVA: A Nonvariational Code for Solving the MHD Stability of Axisymmetric Toroidal Plasmas.” *Journal of Computational Physics* 71 (1): 124–146.
- Cook, Shane. 2012. *CUDA Programming: A Developer's Guide to Parallel Computing with GPUs*. Waltham, MA: Newnes.
- Dagum, Leonardo, and Ramesh Menon. 1998. “OpenMP: An Industry Standard API for Shared-Memory Programming.” *IEEE Computational Science and Engineering* 5 (1): 46–55.
- Duan, Le, Xiaowen Wang, and Xiaolin Zhong. 2010. “A High-Order cut-Cell Method for Numerical Simulation of Hypersonic Boundary-Layer Instability with Surface Roughness.” *Journal of Computational Physics* 229 (19): 7207–7237.
- Evans, T. E., R. A. Moyer, J. G. Watkins, P. R. Thomas, T. H. Osborne, J. A. Boedo, M. E. Fenstermacher, K. H. Finken, R. J. Groebner, and M. Groth. 2005. “Suppression of Large Edge Localized Modes in High Confinement DIII-D Plasmas with a Stochastic Magnetic Boundary.” *Journal of Nuclear Materials* 337: 691–696.
- Farber, Rob. 2016. *Parallel Programming with OpenACC*. Cambridge, MA: Newnes.
- Gupta, Sourendu, and Pushan Majumdar. 2018. “Accelerating Lattice QCD Simulations with 2 Flavors of Staggered Fermions on Multiple GPUs Using OpenACC—A First Attempt.” *Computer Physics Communications* 228: 44–53.
- Kaeli, David R, Perhaad Mistry, Dana Schaa, and Dong Ping Zhang. 2015. *Heterogeneous Computing with OpenCL 2.0*. Waltham, MA: Morgan Kaufmann.
- Kan, Guangyuan, Xiaoyan He, Liuqian Ding, Jiren Li, Yang Hong, Depeng Zuo, Minglei Ren, Tianjie Lei, and Ke Liang. 2018. “Fast Hydrological Model Calibration Based on the Heterogeneous Parallel Computing Accelerated Shuffled Complex Evolution Method.” *Engineering Optimization* 50 (1): 106–119.
- Kan, Guangyuan, Xiaoyan He, Liuqian Ding, Jiren Li, Ke Liang, and Yang Hong. 2017. “A Heterogeneous Computing Accelerated SCE-UA Global Optimization Method Using OpenMP, OpenCL, CUDA, and OpenACC.” *Water Science and Technology*: wst2017322 76: 1640–1651.
- Kan, Guangyuan, Xiaoyan He, Jiren Li, Liuqian Ding, Yang Hong, Hongbin Zhang, Ke Liang, and Mengjie Zhang. 2019. “Computer Aided Numerical Methods for Hydrological

- Model Calibration: An Overview and Recent Development.” *Archives of Computational Methods in Engineering* 26: 35–59.
- Kan, Guangyuan, Tianjie Lei, Ke Liang, Jiren Li, Liuqian Ding, Xiaoyan He, Haijun Yu, Dawei Zhang, Depeng Zuo, and Zhenxin Bao. 2017. “A Multi-Core CPU and Many-Core GPU Based Fast Parallel Shuffled Complex Evolution Global Optimization Approach.” *IEEE Transactions on Parallel and Distributed Systems* 28 (2): 332–344.
- Kan, Guangyuan, Ke Liang, Jiren Li, Liuqian Ding, Xiaoyan He, Youbing Hu, and Mark Amo-Boateng. 2016. “Accelerating the SCE-UA Global Optimization Method Based on Multi-Core CPU and Many-Core GPU.” *Advances in Meteorology* 2016.
- Kan, Guangyuan, Mengjie Zhang, Ke Liang, Hao Wang, Yunzhong Jiang, Jiren Li, Liuqian Ding, Xiaoyan He, Yang Hong, and Depeng Zuo. 2018. “Improving Water Quantity Simulation & Forecasting to Solve the Energy-Water-Food Nexus Issue by Using Heterogeneous Computing Accelerated Global Optimization Method.” *Applied Energy* 210: 420–433.
- Khodak, A. 2015. “Adaptation of General Purpose CFD Code for Fusion MHD Applications.” Paper presented at the 2015 IEEE 26th Symposium on Fusion Engineering (SOFE), 31 May–4 June 2015.
- Kraus, Jiri, Michael Schlottke, Andrew Adinets, and Dirk Pleiter. 2014. “Accelerating a C++ CFD Code with OpenACC.” Proceedings of the First Workshop on Accelerator Programming Using Directives, 47–54. New Orleans, Louisiana: IEEE Press.
- Lao, L. L., H. St John, R. D. Stambaugh, A. G. Kellman, and W. Pfeiffer. 1985. “Reconstruction of Current Profile Parameters and Plasma Shapes in Tokamaks.” *Nuclear Fusion* 25 (11): 1611.
- Liang, Y., H. R. Koslowski, P. R. Thomas, E. Nardon, B. Alper, P. Andrew, Y. Andrew, G. Arnoux, Y. Baranov, and M. Bécoulet. 2007. “Active Control of Type-I Edge-Localized Modes with $n = 1$ Perturbation Fields in the JET Tokamak.” *Physical Review Letters* 98 (26): 265004.
- Markidis, S., J. Gong, M. Schliephake, E. Laure, A. Hart, D. Henty, K. Heisey, and P. Fischer. 2015. “OpenACC Acceleration of the Nek5000 Spectral Element Code.” *International Journal of High Performance Computing Applications* 29 (3): 311–319. doi:10.1177/1094342015576846.
- Mukherjee, Rupak, Rajaraman Ganesh, Vinod Saini, Udaya Maurya, Nagavijayalakshmi Vydyanathan, and Bharatkumar Sharma. 2018. “Three Dimensional Pseudo-Spectral Compressible Magnetohydrodynamic GPU Code for Astrophysical Plasma Simulation.” arXiv preprint arXiv:1810.12707.
- Otten, Matthew, Jing Gong, Azamat Mametjanov, Aaron Vose, John Levesque, Paul Fischer, and Misun Min. 2016. “An MPI/OpenACC Implementation of a High-Order Electromagnetics Solver with GPUDirect Communication.” *The International Journal of High Performance Computing Applications* 30 (3): 320–334.
- Peng, Ivy Bo, Stefano Markidis, Andris Vaivads, Juris Vencels, Jan Deca, Giovanni Lapenta, Alistair Hart, and Erwin Laure. 2015. “Acceleration of a Particle-in-Cell Code for Space Plasma Simulations with OpenACC.” Paper presented at the EGU General Assembly Conference Abstracts.
- Renner, H., J. Boscary, V. Erckmann, H. Greuner, H. Grote, J. Sapper, E. Speth, F. Wesner, and M. Wanner. 2000. “The Capabilities of Steady State Operation at the Stellarator W7-X with Emphasis on Divertor Design.” *Nuclear Fusion* 40 (6): 1083.
- Song, Yun Tao, Song Tao Wu, Jian Gang Li, Bao Nian Wan, Yuan Xi Wan, Peng Fu, Min You Ye, Jin Xing Zheng, Kun Lu, and Xianggao Gao. 2014. “Concept Design of CFETR Tokamak Machine.” *IEEE Transactions on Plasma Science* 42 (3): 503–509.
- Walker, David W, and Jack J Dongarra. 1996. “MPI: A Standard Message Passing Interface.” *Supercomputer* 12: 56–68.
- Wan, Baonian. 2009. “Recent Experiments in the EAST and HT-7 Superconducting Tokamaks.” *Nuclear Fusion* 49 (10): 104011.
- Wang, Yichao, James Lin, Linjin Cai, William Tang, Stephane Ethier, Bei Wang, Simon See, and Satoshi Matsuoka. 2016. “Porting and Optimizing gtc-p on Taihulight Supercomputer with Sunway Openacc.” Paper presented at HPC China, Xi’an, China, October 27. <https://hpc.sjtu.edu.cn/slides.pdf>.
- Wang, S., and Z. W. Ma. 2015. “Influence of Toroidal Rotation on Resistive Tearing Modes in Tokamaks.” *Physics of Plasmas* 22 (12): 122504.
- Wang, S., Z. W. Ma, and Wei Zhang. 2016. “Influence of Driven Current on Resistive Tearing Mode in Tokamaks.” *Physics of Plasmas* 23 (5): 052503.
- Wang, Endong, Qing Zhang, Bo Shen, Guangyong Zhang, Xiaowei Lu, Qing Wu, and Yajuan Wang. 2014. “Intel Math Kernel Library.” In *High-Performance Computing on the Intel® Xeon Phi™*, 167–188. Beijing, China: Springer.
- Zhang, W., Z. W. Ma, and S. Wang. 2017. “Hall Effect on Tearing Mode Instabilities in Tokamak.” *Physics of Plasmas* 24 (10): 102510.
- Zhang, W., S. Wang, and Z. W. Ma. 2017. “Influence of Helical External Driven Current on Nonlinear Resistive Tearing Mode Evolution and Saturation in Tokamaks.” *Physics of Plasmas* 24 (6): 062510.
- Zhu, J., Z. W. Ma, and S. Wang. 2016. “Hybrid Simulations of Alfvén Modes Driven by Energetic Particles.” *Physics of Plasmas* 23 (12): 122506.
- Zhu, Jia, Zhiwei Ma, Sheng Wang, and Wei Zhang. 2018. “Non-linear Dynamics of Toroidal Alfvén Eigenmodes in the Presence of Tearing Modes.” *Nuclear Fusion* 58.